

1. a) $\frac{1}{n} \sum_{i=1}^n \nabla_{\theta} f_{\theta}(x_i)$ estimates $E[\nabla_{\theta} f_{\theta}(x)]$, which is only part of $\nabla_{\theta} L(\theta)$, thus making the result biased towards the other term.

b) We simply average the n samples: $\frac{1}{n} \sum_{i=1}^n f_{\theta}(\mu + \sigma^{1/2} \epsilon)$